

Understanding the msghdr structure from sys/socket.h

Ask Question



6



3

I'm trying to understand the following members of the `msg_hdr` structure of the [sys/socket.h](#) lib.

- `struct iovec *msg_iov` scatter/gather array
- `void *msg_control` ancillary data, see below

It states below:

Ancillary data consists of a sequence of pairs, each consisting of a `cmsg_hdr` structure followed by a data array. The data array contains the ancillary data message, and the `cmsg_hdr` structure contains descriptive information that allows an application to correctly parse the data.

I'm assuming the `msg_hdr` struct, contains the protocol-header information? if so... `*msg_iov` is the input/output "vector" of parameters in the request/response? and the `*msg_control` contains the response messages?

`c` `linux` `sockets` `network-programming`

share improve this question

add a comment

edited Apr 27 '18 at 12:40



Mateusz Piotrowski
3,880 ● 6 ● 30 ● 51

asked Sep 15 '15 at 19:04



Jordan Davis
649 ● 11 ● 30

1 Answer

active oldest votes



`msg_iov` is an array of input/output buffers with length `msg_iovlen`. Each member of this array contains a pointer to a data buffer and the size of the buffer. This is where the data to read/write lives. It allows you to read/write to an array of buffers which are not necessarily in contiguous memory regions.



`msg_control` points to a buffer of size `msg_controllen` that contains additional information about the packet. To read this field, you first need to declare a `struct cmsghdr *` (let's call it `cmhdr`). You populate this by calling `MSG_FIRSTHDR()` the first time, passing it the address of the `msg_hdr` struct, and `MSG_NXTHDR()` each subsequent time, passing it the address of the `msg_hdr` struct and the current value of `cmhdr`.

From the `msg_control`, you can find interesting things like the destination IP of the packet (useful for multicast) and the contents of the TOS/DSCP byte in the IP header (useful for custom congestion control protocols), among others. In most cases, you'll need to make a `setsockopt` call to enable receiving this data. In the examples given, the `IP_PKTINFO` and `IP_TOS` options need to be enabled.

See the [cmsghdr\(3\) manpage](#) for more details.

The source IP and port, are not in `msg_control`, but are in `msg_name` which expects a pointer to a `struct sockaddr` with length `msg_namelen`.

Here's an example of how to use this:

```
struct msg_hdr mhdr;
struct iovec iov[1];
struct cmsghdr *cmhdr;
char control[1000];
struct sockaddr_in sin;
char databuf[1500];
unsigned char tos;

mhdr.msg_name = &sin;
mhdr.msg_namelen = sizeof(sin);
mhdr.msg_iov = iov;
mhdr.msg_iovlen = 1;
mhdr.msg_control = &control;
mhdr.msg_controllen = sizeof(control);
```

```

iov[0].iov_base = databuf;
iov[0].iov_len = sizeof(databuf);
memset(databuf, 0, sizeof(databuf));
if ((*len = recvmsg(sock, &mhdr, 0)) == -1) {
    perror("error on recvmsg");
    exit(1);
} else {
    cmhdr = CMSG_FIRSTHDR(&mhdr);
    while (cmhdr) {
        if (cmhdr->cmsg_level == IPPROTO_IP && cmhdr->cmsg_type == IP_TOS) {
            // read the TOS byte in the IP header
            tos = ((unsigned char *)CMSG_DATA(cmhdr))[0];
        }
        cmhdr = CMSG_NXTHDR(&mhdr, cmhdr);
    }
    printf("data read: %s, tos byte = %02X\n", databuf, tos);
}

```

share improve this answer

edited Sep 15 '15 at 19:35

answered Sep 15 '15 at 19:27



dbush

97.1k ● 13 ● 104 ● 138

I got it now, thanks again dbush! – [Jordan Davis](#) Sep 15 '15 at 19:48

why is `struct sockaddr_in` ? shouldn't it be `struct sockaddr` ? – [Jordan Davis](#) Sep 15 '15 at 21:39

1 `struct sockaddr` is a generic type for getting socket info. For an IPv4 socket, the specific type is `struct sockaddr_in`. For a UNIX domain socket, it's `struct sockaddr_un`, and for an IPv6 socket, it's `struct sockaddr_in6` – [dbush](#) Sep 15 '15 at 21:42

Haha okay. Where the hell do you see that though lol? I don't even see in the spec at all. – [Jordan Davis](#) Sep 15 '15 at 21:45

1 `man 7 ip`, `man 7 unix`, and `man 7 ipv6`, as well as `man 7 socket`. – [dbush](#) Sep 15 '15 at 21:49

[show 3 more comments](#)

Your Answer

B

I



Sign up or [log in](#)

 Sign up using Google

 Sign up using Facebook

 Sign up using Email and Password

Post as a guest

Name

Email

Required, but never shown

Post Your Answer

By clicking "Post Your Answer", you acknowledge that you have read our updated [terms of service](#), [privacy policy](#) and [cookie policy](#), and that your continued use of the website is subject to these policies.

Not the answer you're looking for? Browse other questions tagged [c](#) [linux](#) [sockets](#) [network-programming](#) or [ask your own question](#).

asked 3 years, 4 months ago

viewed 5,529 times

active 9 months ago

BLOG



How the 2019 Stack Overflow Developer Survey Came to Be (And Your Last Chance...

Linked

0

[sockaddr structure - \(sys/socket.h\)](#)

Related

6

[In C: How to set a pointer to a structure member that is an array?](#)

209

[Purpose of Unions in C and C++](#)











5
















[Locate unused structures and structure-](#)

members

- 2 How to construct and manage this complex C data structure?
- 0 Can ancillary data be portably allocated?
- 2 Linux group and user management
- 4 Data Type - socklen_t, sa_family_t
- 0 sockaddr structure - (sys/socket.h)
- 1 Socket Programming (sys/socket.h)
- 3 Using sendmsg/sendmmsg with raw ethernet frames

Hot Network Questions

-  Russian Dolls - How do they reproduce?
-  Heisenberg groups, Carnot groups and contact forms
-  Sigma notation for non-integer steps
-  Determinantal symmetry: proof requested
-  What's a mean field variational family?
-  A 4 x 4 Magic Square with Pairwise Relatively Prime Entries
-  Novel with science masqueraded as religion; rebels use the occult to fight for freedom
-  Is there a loss of entropy by hashing an N-bit random key to produce an N-bit hash?
-  Is it practical to make English the global language?
-  Is it bad to have no gender variety?

-  What are the alternatives for Xposed framework on Pie? Preferably without triggering Safety Net?
-  Why does Python allow out-of-range slice indexes for sequences?
-  N-movers: How much of the infinite board can I reach?
-  Can modify memory be used to change a revenant's target?
-  A man killed my son and married one of the witnesses. Can the witness's statements still be used in court against him?
-  How to retain a healthy relationship with colleagues after beating them at golf?
-  Is it inconvenient for the DM if I make a pet to go with my character?
-  How to deal with a "protest" assignment?
-  Why is plagiarising words such a big deal?
-  Is it appropriate to request that the department hire a female faculty member?
-  Is there an EU law that prevents repair of electrical devices?
-  find value of trig integral
-  What was the first games console to support linking consoles together for local multiplayer?
-  Am I responsible for finding my own replacement?
-  question feed



STACK OVERFLOW

Questions

PRODUCTS

Teams

COMPANY

About

STACK EXCHANGE
NETWORK

[Blog](#) [Facebook](#) [Twitter](#) [LinkedIn](#)

[Jobs](#)

[Developer Jobs Directory](#)

[Salary Calculator](#)

[Help](#)

[Mobile](#)

[Disable Responsiveness](#)

[Talent](#)

[Engagement](#)

[Enterprise](#)

[Press](#)

[Work Here](#)

[Legal](#)

[Privacy Policy](#)

[Contact Us](#)

[Technology](#) >

[Life / Arts](#) >

[Culture / Recreation](#) >

[Science](#) >

[Other](#) >

site design / logo © 2019 Stack Exchange Inc; user
contributions licensed under cc by-sa 3.0 with
attribution required. rev 2019.2.8.32848